



人工智能实践

Artificial Intelligence Practice

DCS3015 Autumn 2022

Chao Yu (余超)

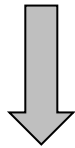
School of Computer Science and Engineering
Sun Yat-Sen University

Lecture 8: Model-based RL

15th December 2022

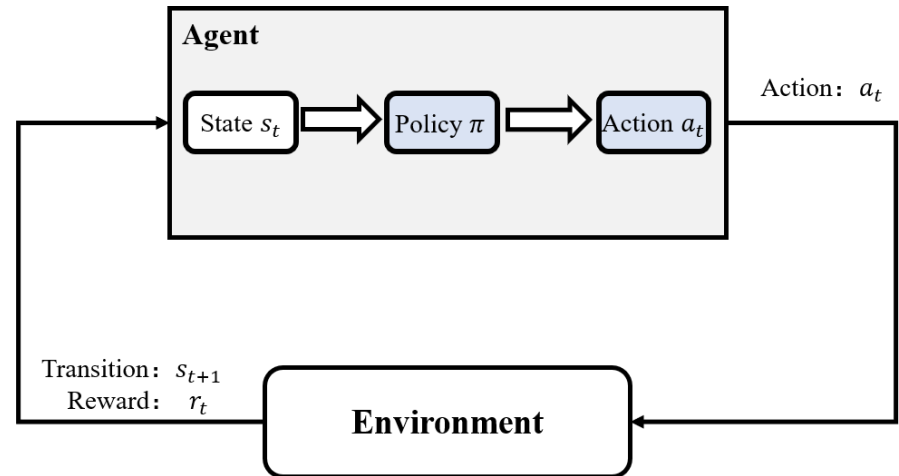
Quick Recap

- States $S \in \mathbb{R}^s$
- Actions $A \in \mathbb{R}^a$
- Reward Function $R: S \times A \rightarrow \mathbb{R}$
- Transition Function $T: S \times A \rightarrow S$
- Discounted Factor $\gamma \in (0,1)$
- Policy $\pi: S \rightarrow A$



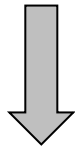
$$\max_{\pi} \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t)$$

Subject to $a_t = \pi(s_t), s_{t+1} = T(s_t, a_t)$



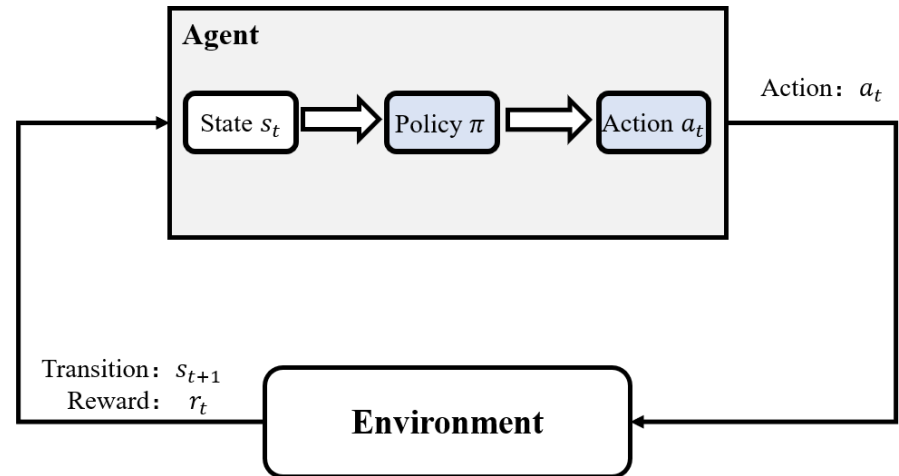
Quick Recap

- States $S \in \mathbb{R}^s$
- Actions $A \in \mathbb{R}^a$
- Reward Function $R: S \times A \rightarrow \mathbb{R}$
- Transition Function $T: S \times A \rightarrow S$
- Discounted Factor $\gamma \in (0,1)$
- Policy $\pi: S \rightarrow A$



$$\max_{\pi} \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t)$$

Subject to $a_t = \pi(s_t)$, $s_{t+1} = T(s_t, a_t)$





Model-free RL vs. Model-based RL

Collect data

$$\mathcal{D} = \{s_t, a_t, r_{t+1}, s_{t+1}\}_{t=0}^T$$

Model-free: learn policy directly from data

$$\mathcal{D} \rightarrow \pi \quad \text{e.g. } Q\text{-learning, policy gradient}$$

Model-based: learn model, then use it to learn or improve a policy

$$\mathcal{D} \rightarrow f \rightarrow \pi$$

Model-free RL vs. Model-based RL

Collect data

$$\mathcal{D} = \{s_t, a_t, r_{t+1}, s_{t+1}\}_{t=0}^T$$

Model-free: learn policy directly from data

$$\mathcal{D} \rightarrow \pi \quad \text{e.g. } Q\text{-learning, policy gradient}$$

Model-based: learn model, then use it to learn or improve a policy

$$\mathcal{D} \rightarrow f \rightarrow \pi$$

□ Why Model-based ?

	Model-free	Model-based
Asymptotic Performance	+	- / +
Computation	+	-
Sample Efficiency	-	+
Exploration	-	+

Model-free RL vs. Model-based RL

Collect data

$$\mathcal{D} = \{s_t, a_t, r_{t+1}, s_{t+1}\}_{t=0}^T$$

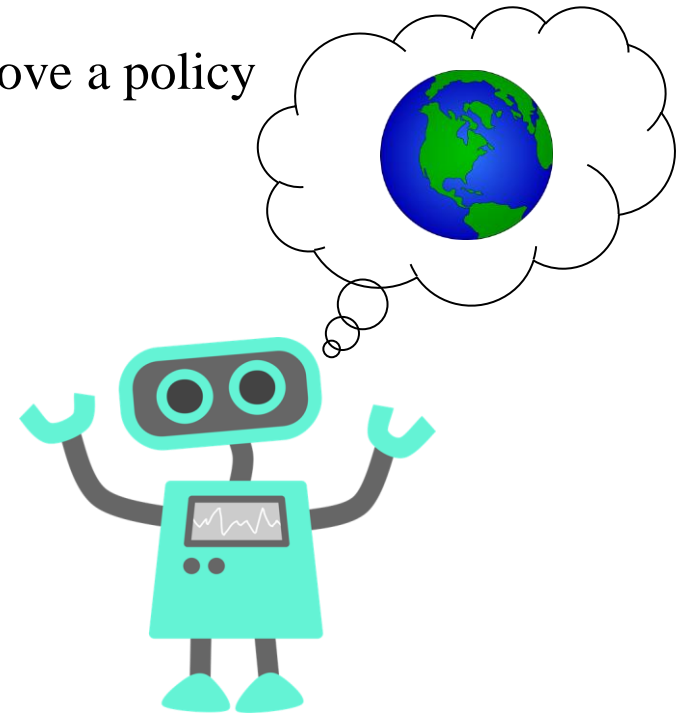
Model-free: learn policy directly from data

$$\mathcal{D} \rightarrow \pi \quad \text{e.g. } Q\text{-learning, policy gradient}$$

Model-based: learn model, then use it to learn or improve a policy

$$\mathcal{D} \rightarrow f \rightarrow \pi$$

□ **What is a Model ?**



What is a Model ?

Collect data

$$\mathcal{D} = \{s_t, a_t, r_{t+1}, s_{t+1}\}_{t=0}^T$$

Model-free: learn policy directly from data

$$\mathcal{D} \rightarrow \pi \quad \text{e.g. } Q\text{-learning, policy gradient}$$

Model-based: learn model, then use it to learn or improve a policy

$$\mathcal{D} \rightarrow f \rightarrow \pi$$

□ What is a Model ?

Definition: a model is a representation that **explicitly** encodes knowledge about the structure of the environment and task. `state_next, reward, terminate, infos = env.step(action)`

A transition/dynamics model:

$$s_{t+1} = f_s(s_t, a_t)$$

A rewards model:

$$r_{t+1} = f_r(s_t, a_t)$$

A inverse transition/dynamics model:

$$a_t = f_s^{-1}(s_t, s_{t+1})$$

What is a Model ?

Collect data

$$\mathcal{D} = \{s_t, a_t, r_{t+1}, s_{t+1}\}_{t=0}^T$$

Model-free: learn policy directly from data

$$\mathcal{D} \rightarrow \pi \quad \text{e.g. } Q\text{-learning, policy gradient}$$

Model-based: learn model, then use it to learn or improve a policy

$$\mathcal{D} \rightarrow f \rightarrow \pi$$

□ What is a Model ?

Definition: a model is a representation that **explicitly** encodes knowledge about the structure of the environment and task. `state_next, reward, terminate, infos = env.step(action)`

A transition/dynamics model:

$$s_{t+1} = f_s(s_t, a_t)$$

A rewards model:

$$r_{t+1} = f_r(s_t, a_t)$$

A inverse transition/dynamics model:

$$a_t = f_s^{-1}(s_t, s_{t+1})$$

Typically what is meant by the model in model-based RL

What is a Model ?

Collect data

$$\mathcal{D} = \{s_t, a_t, r_{t+1}, s_{t+1}\}_{t=0}^T$$

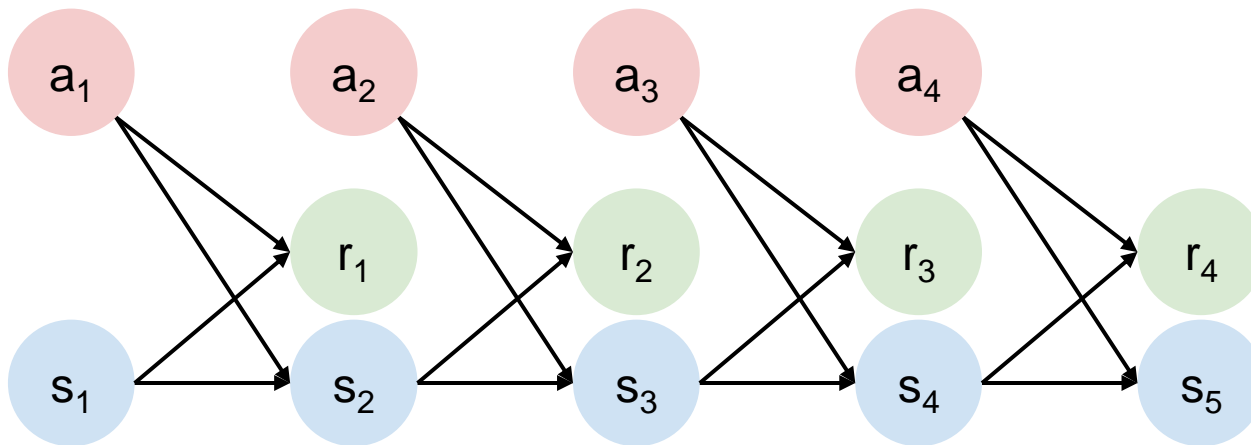
Model-free: learn policy directly from data

$$\mathcal{D} \rightarrow \pi \quad \text{e.g. } Q\text{-learning, policy gradient}$$

Model-based: learn model, then use it to learn or improve a policy

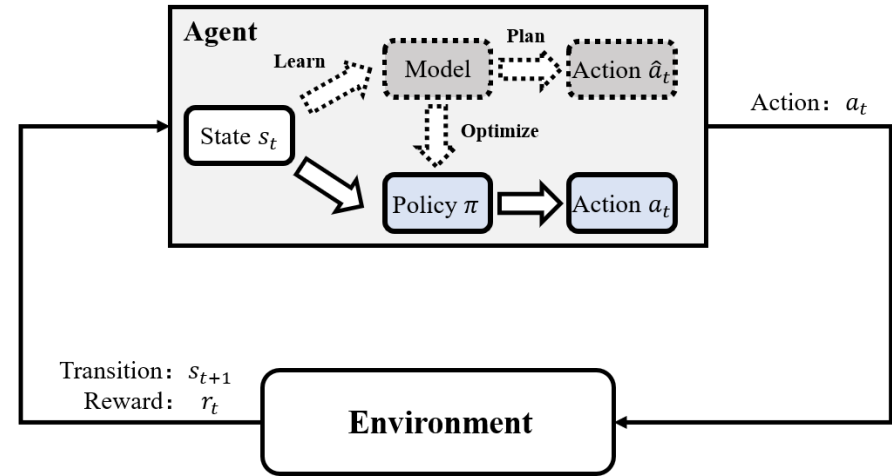
$$\mathcal{D} \rightarrow f \rightarrow \pi$$

□ What is a Model ?



Model-based RL

□ Model-based RL Process:



Run base policy $\pi_0(a_t|s_t)$ to collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

For _ do

Learn transition/dynamics model $f(s, a)$

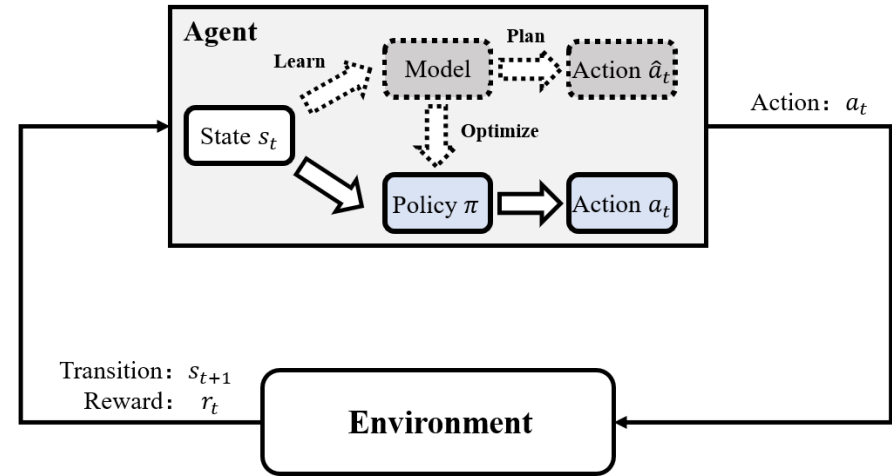
Use model $f(s, a)$ to assist agent to update policy $\pi(a_t|s_t)$

Execute $\pi(a_t|s_t)$ and add the resulting data $\{(s_t, a_t, s_{t+1})_j\}$ to \mathcal{D}

End For

Model-based RL

□ Model-based RL Process:



Run base policy $\pi_0(a_t|s_t)$ to collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

For _ do

Learn transition/dynamics model $f(s, a)$

Model Building

Use model $f(s, a)$ to assist agent to update policy $\pi(a_t|s_t)$

Model Utilizing

Execute $\pi(a_t|s_t)$ and add the resulting data $\{(s_t, a_t, s_{t+1})_j\}$ to \mathcal{D}

End For

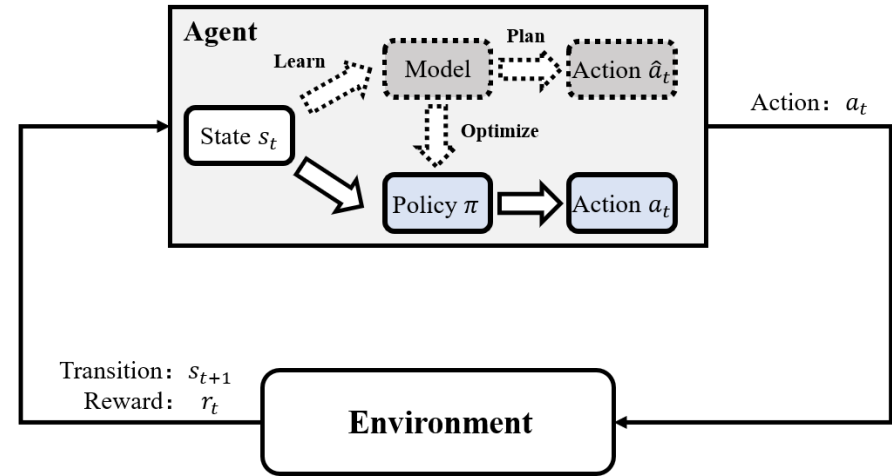
Model-based RL

□ Model Building

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$



Deterministic Model or Probabilistic Model



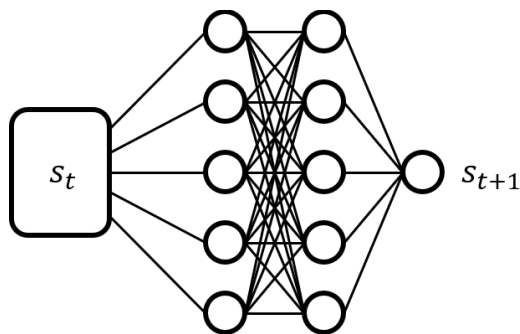
Model-based RL

Model Building

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$



Deterministic Model or Probabilistic Model



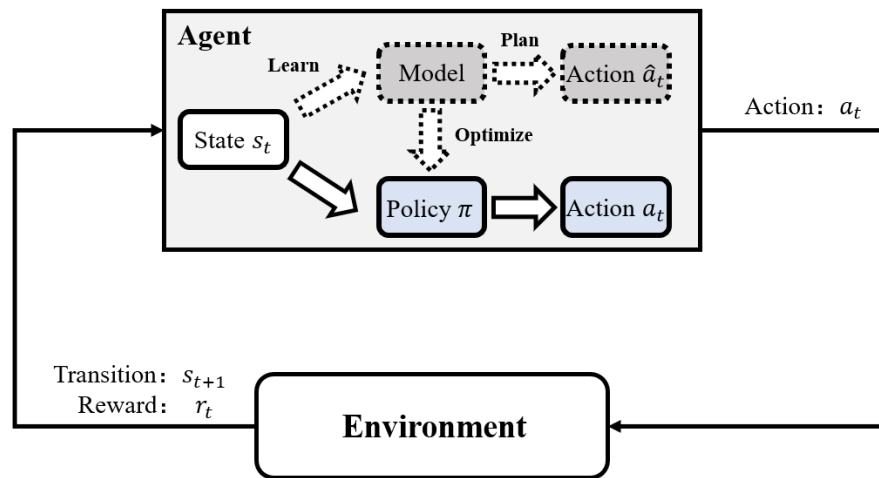
Learn transition model $f(s, a)$ to minimize **Mean Squared Error (MSE)**:

$$\mathcal{L}_{MSE} = \mathbb{E}_{(s_t, a_t, s_{t+1}) \sim \mathcal{D}} [f(s_t, a_t) - s_{t+1}]^2.$$

$$\hat{s}_{t+1} \sim f(s_t, a_t)$$

Incremental Predict:

$$\hat{s}_{t+1} \sim s_t + f(s_t, a_t)$$



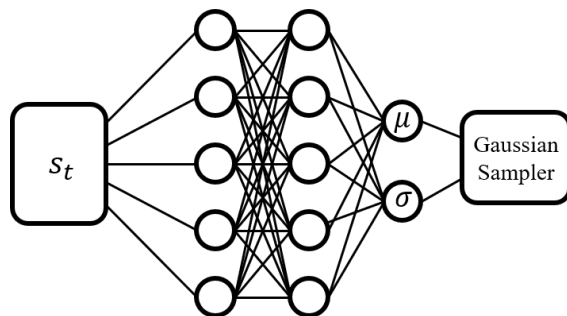
Model-based RL

Model Building

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$



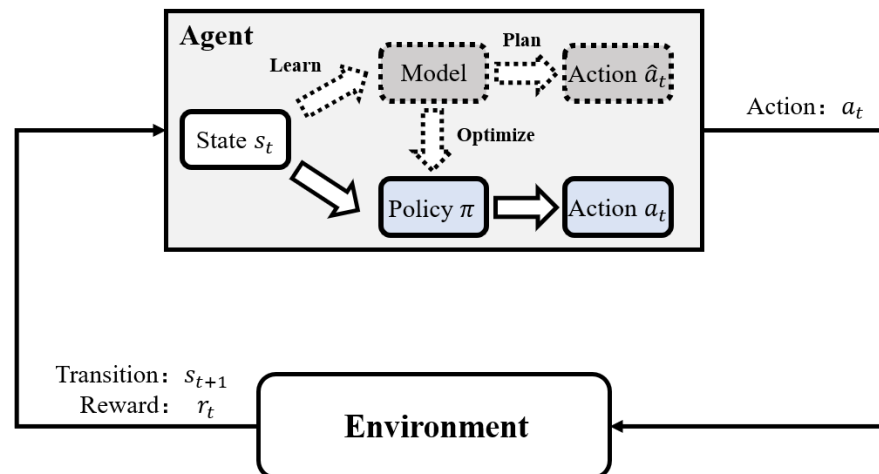
Deterministic Model or **Probabilistic Model**



Learn transition model $f(s, a)$ to optimize **Negative Log Likelihood (NLL)**

$$\mathcal{L}_{NLL} = \mathbb{E}_{(s_t, a_t, s_{t+1}) \sim \mathcal{D}} \left(\frac{[\mu(s_t, a_t) - s_{t+1}]^2}{\sigma(s_t, a_t)} + \log \sigma(s_t, a_t) \right)$$

$$\hat{s}_{t+1} \sim \mathcal{N}(\mu(s_t, a_t), \sigma(s_t, a_t))$$



Model-based RL

Model Building:

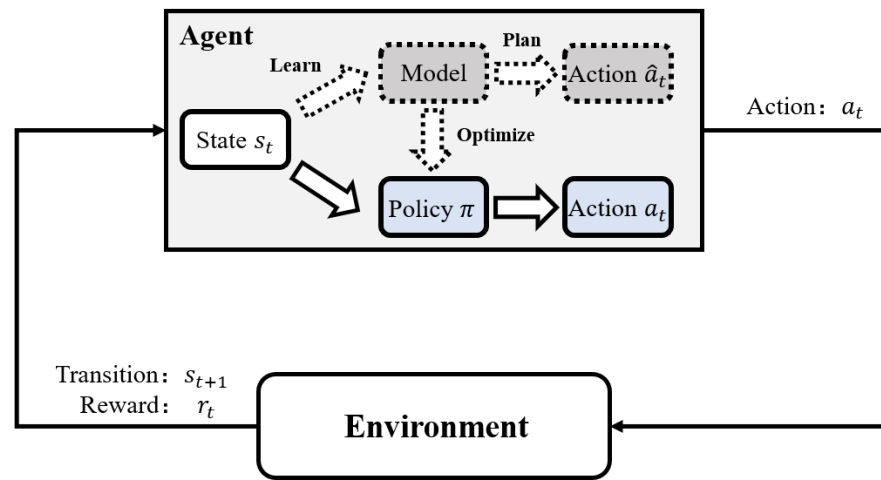
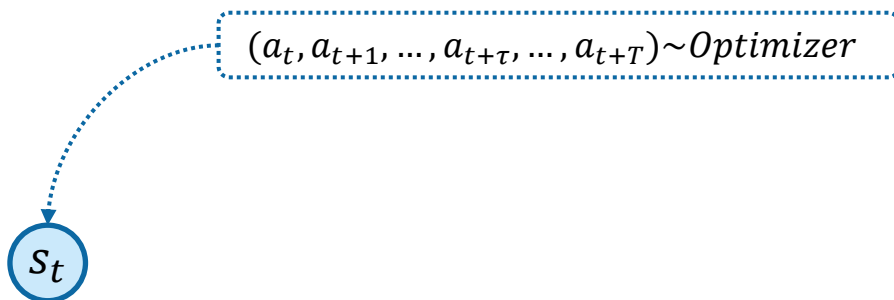
Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$

Decision-Time Method (Plan):

Model Predictive Control (MPC) without Policy Function π



Model-based RL

□ Model Building:

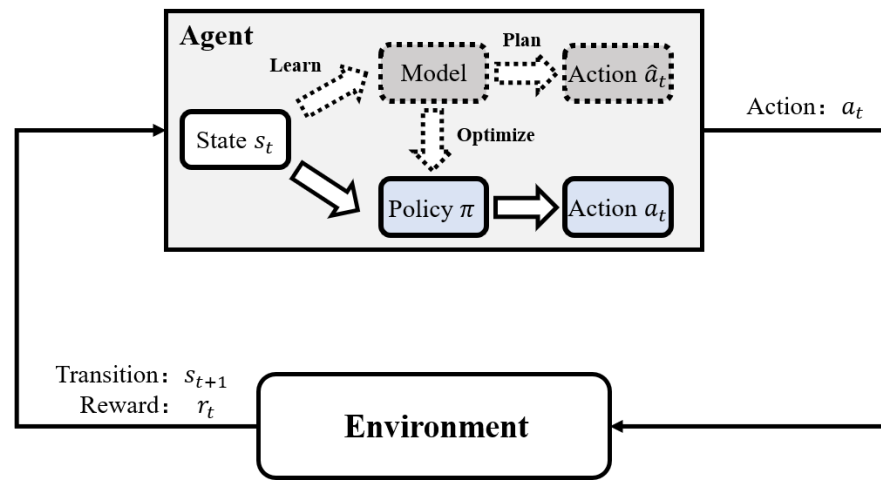
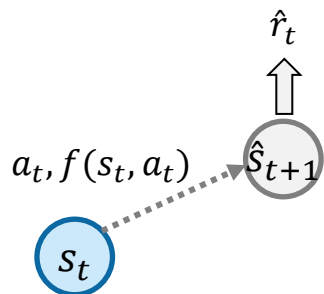
Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$

□ Decision-Time Method (Plan):

Model Predictive Control (MPC) without Policy Function π



Model-based RL

□ Model Building:

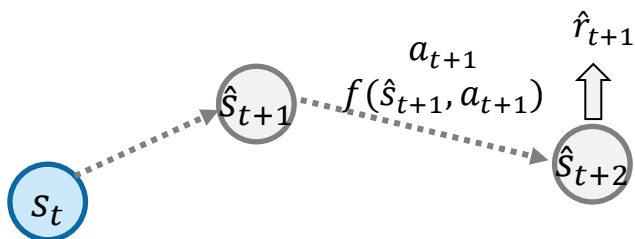
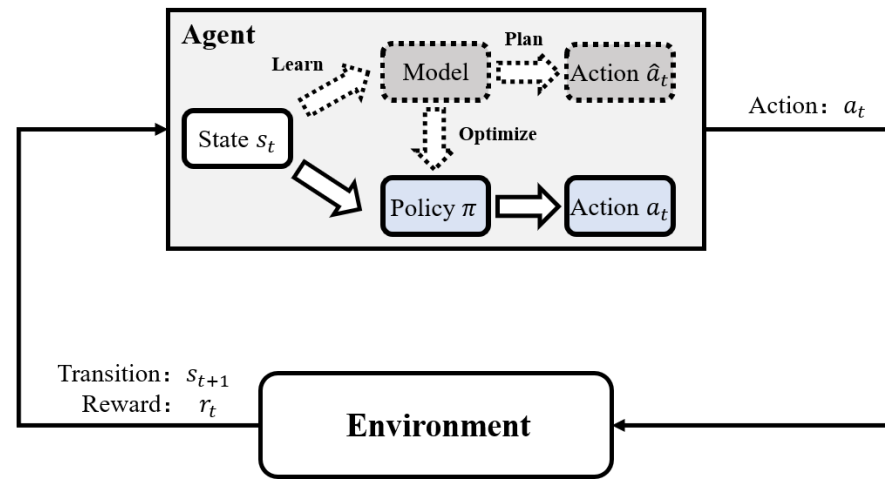
Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$

□ Decision-Time Method (Plan):

Model Predictive Control (MPC) without Policy Function π



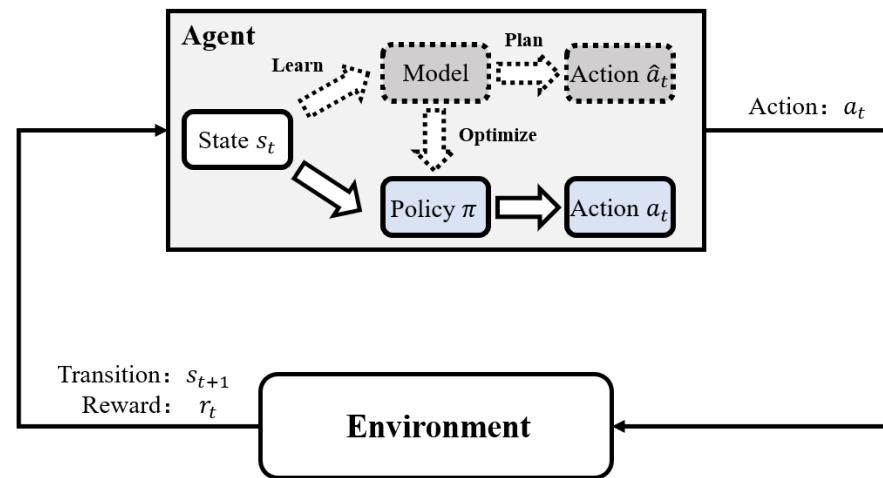
Model-based RL

Model Building:

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

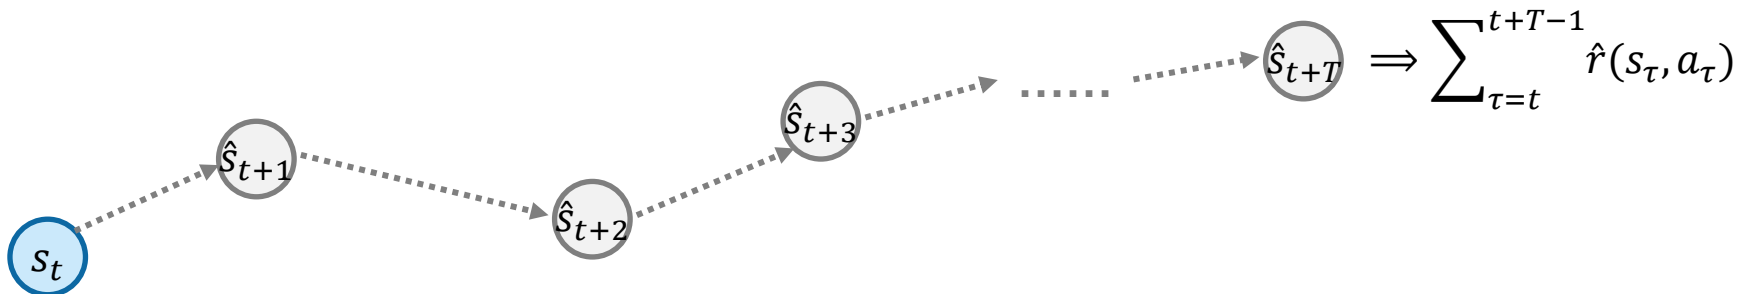
Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$



Decision-Time Method (Plan):

Model Predictive Control (MPC) without Policy Function π



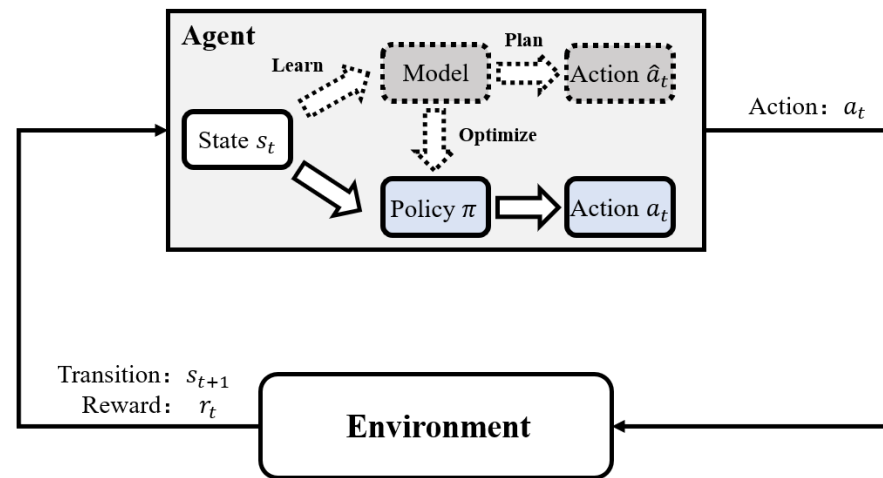
Model-based RL

Model Building:

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

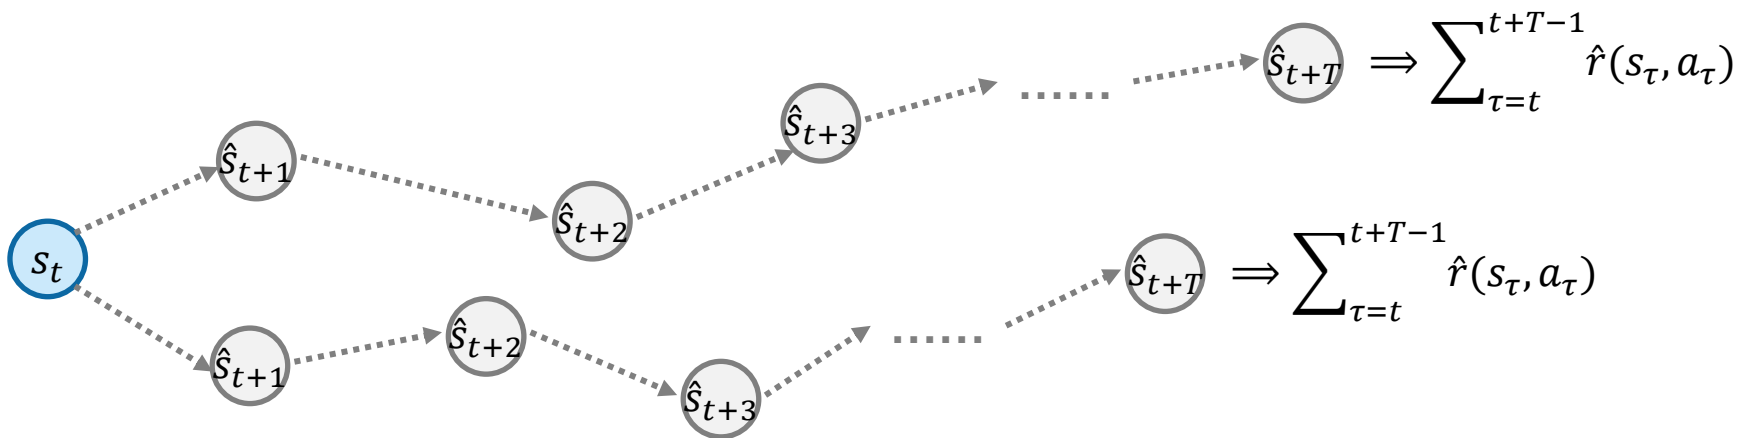
Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$



Decision-Time Method (Plan):

Model Predictive Control (MPC) without Policy Function π



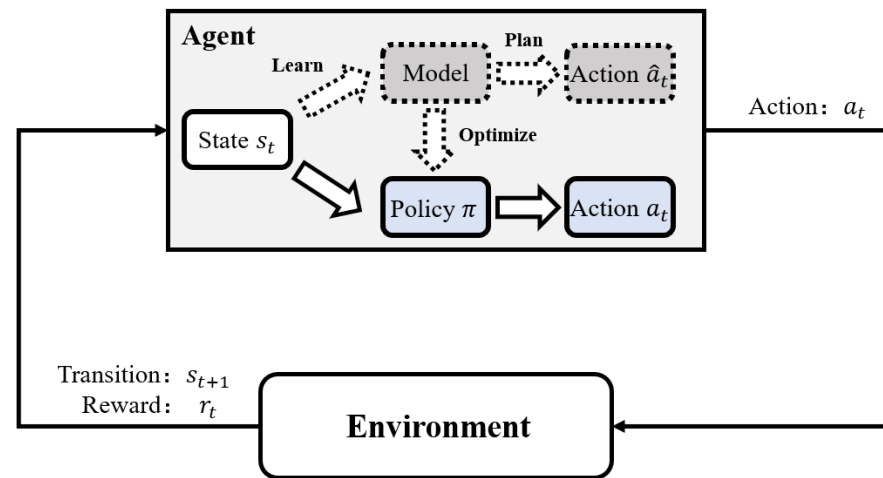
Model-based RL

Model Building:

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

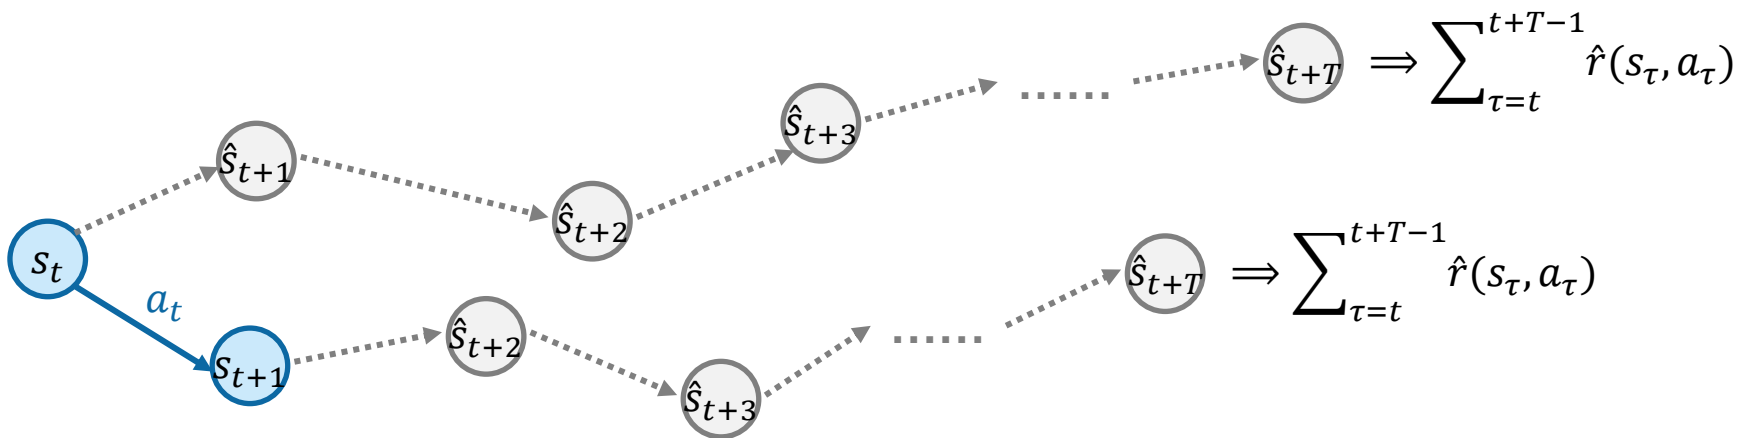
Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$



Decision-Time Method (Plan):

Model Predictive Control (MPC) without Policy Function π



Model-based RL

Model Building:

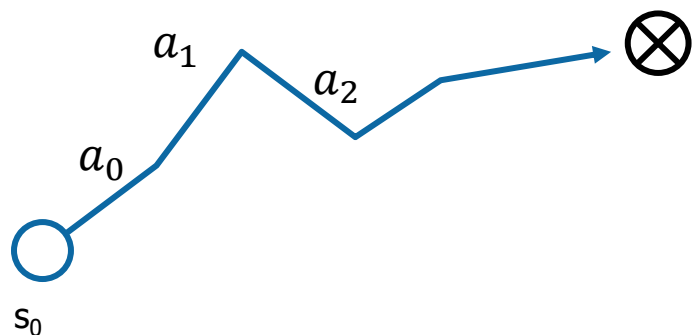
Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$

Decision-Time Method (Plan):

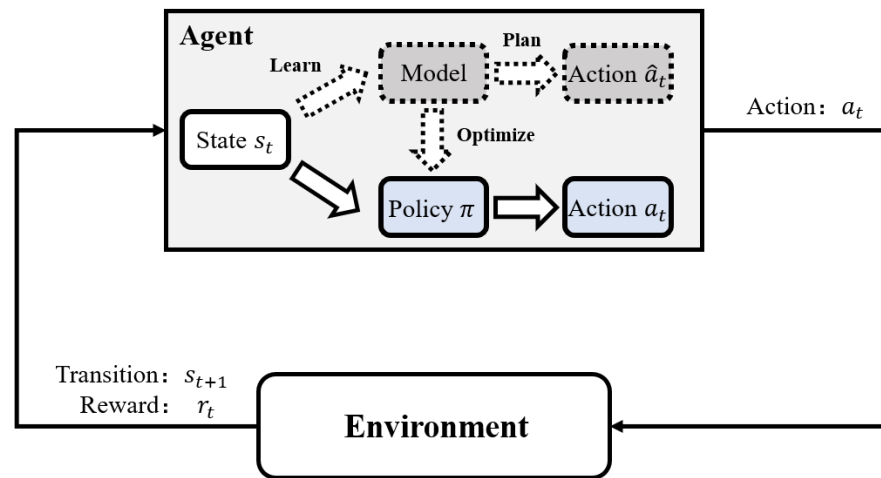
Model Predictive Control (MPC) without Policy Function π



Trajectory Optimization with Action Sequence:

$$s_t \rightarrow a_t \rightarrow s_{t+1} \rightarrow a_{t+1} \rightarrow \dots$$

$$J(a_0, \dots, a_H) = \sum_{t=0}^H \gamma^t r_t$$



Model-based RL

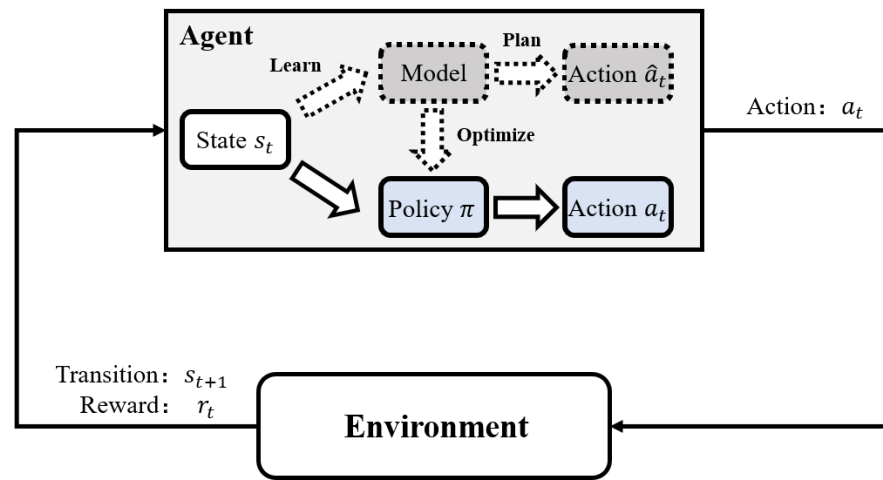
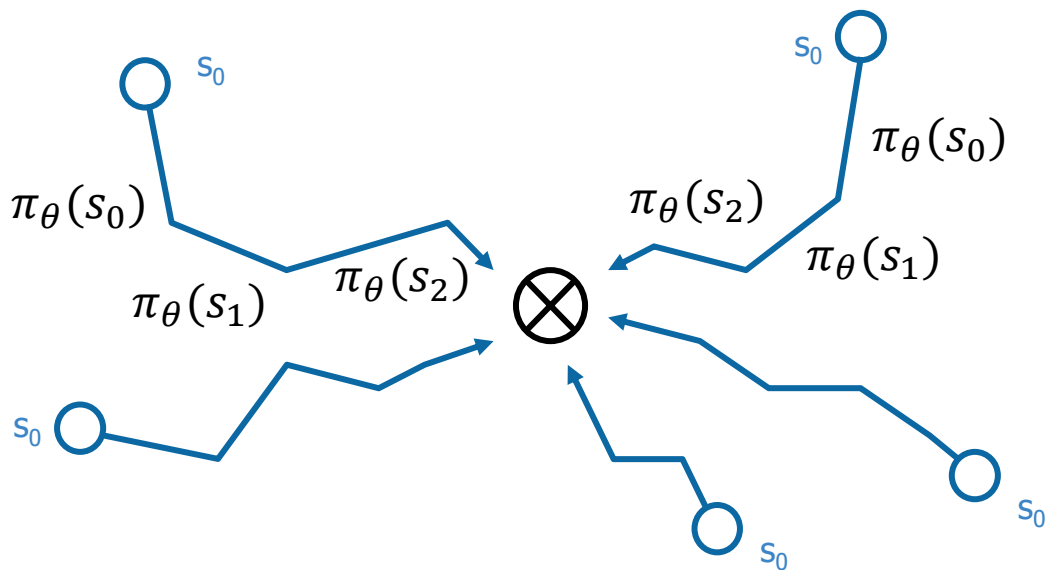
Model Building:

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

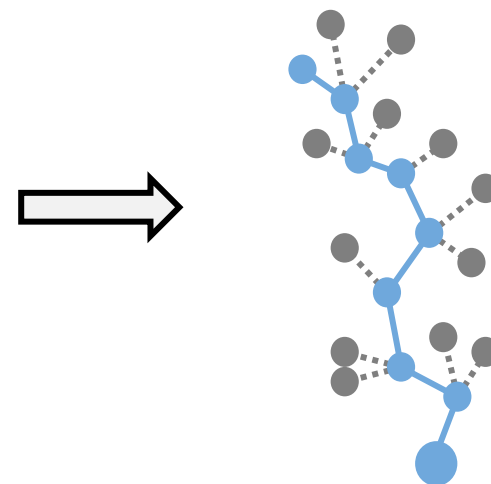
Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$

Background Method (Optimize):



Mix **real** and **model-generated** experience and apply additional policy updates.



Model-based RL

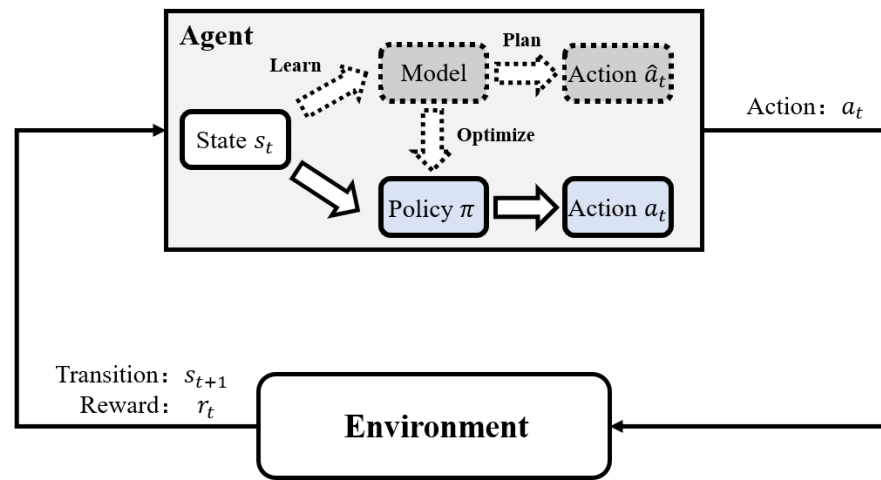
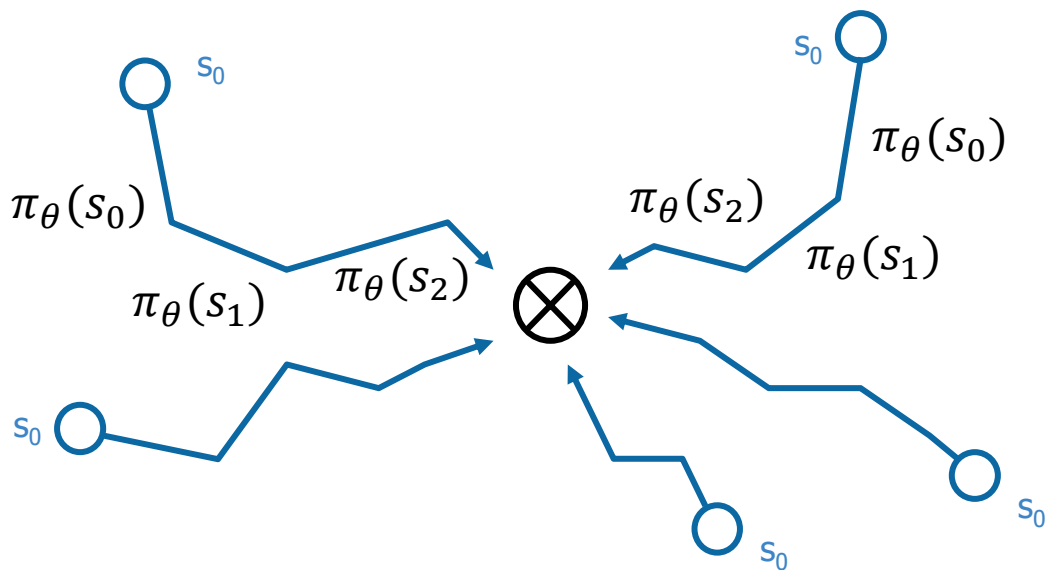
Model Building:

Collect data $\mathcal{D} = \{(s_t, a_t, s_{t+1})_i\}$

Learn transition model $f(s, a)$ to minimize

$$\sum_i \|f(s_t, a_t) - s_{t+1}\|^2$$

Background Method (Optimize):



Policy Optimization with parameter θ :

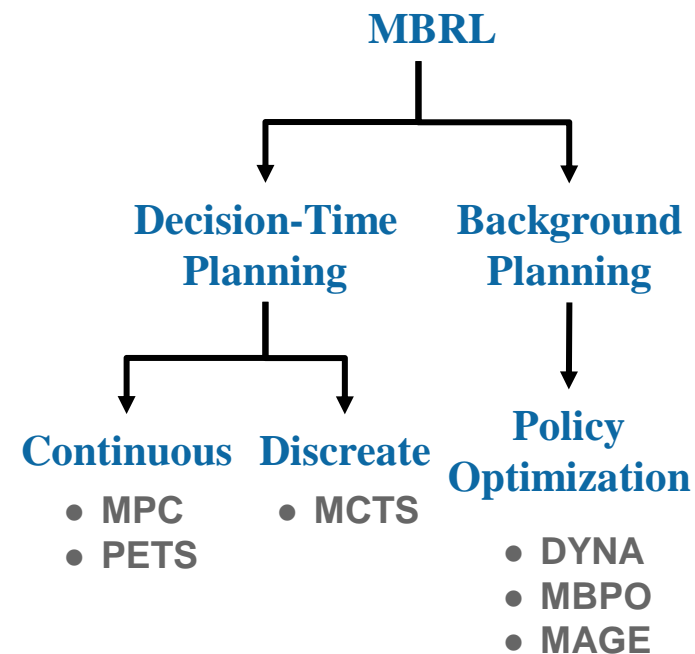
$$s_t \rightarrow \pi_\theta(s_t) \rightarrow s_{t+1} \rightarrow \pi_\theta(s_{t+1}) \rightarrow \dots$$

$$J(\theta) = \mathbb{E}_{s_0} \left[\sum_{t=0}^H \gamma^t r_t \right], \quad a_t = \pi_\theta(s_t)$$

Model-based RL

□ Background vs. Decision-Time

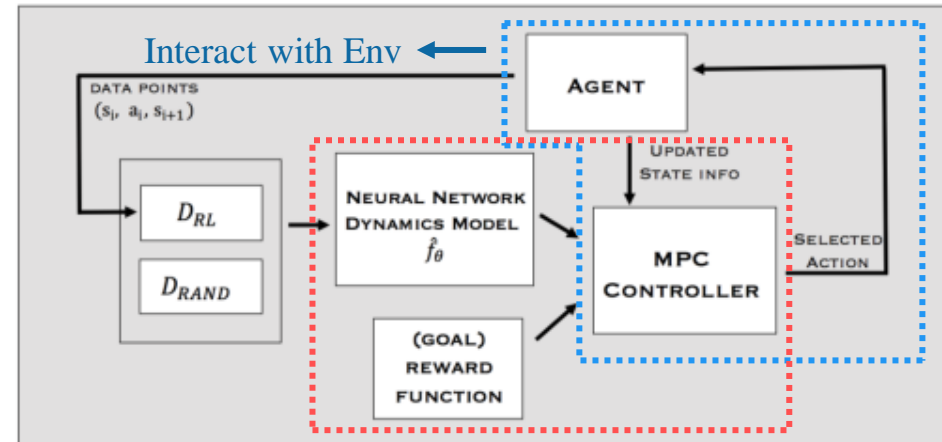
	Background	Decision-Time
Act without Learning	-	+
Unfamiliar situations	-	+
Computation	+	-
Predictability Coherence	+	-
Discrete and Continuous	+	-



Decision-Time

□ MPC with Neural Network Model

- $\mathcal{D}_{RL} = \{(s_t, a_t, s_{t+1})_i\}$ to train neural network dynamic model.
- $\mathcal{D}_{RAND} = \{(a_1, a_2, \dots, a_T)\}$ to store the sampled actions sequence.
- (Goal) Reward Function



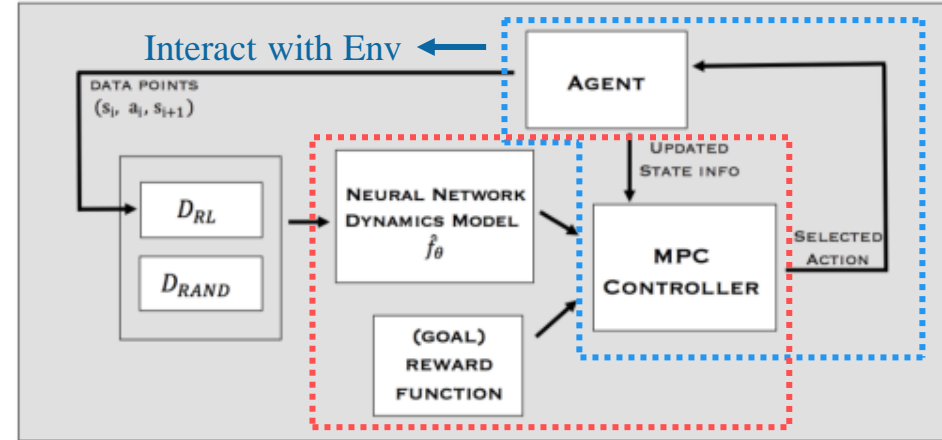
1. Planning based on Model
2. Agent action selection

Decision-Time

□ MPC with Neural Network Model

Algorithm 1 Model-based Reinforcement Learning

- 1: gather dataset $\mathcal{D}_{\text{RAND}}$ of random trajectories
- 2: initialize empty dataset \mathcal{D}_{RL} , and randomly initialize \hat{f}_{θ}
- 3: **for** iter=1 **to** max_iter **do**
- 4: train $\hat{f}_{\theta}(s, a)$ by performing gradient descent on Eqn. 2 using $\mathcal{D}_{\text{RAND}}$ and \mathcal{D}_{RL}
- 5: **for** $t = 1$ **to** T **do**
- 6: get agent's current state s_t
- 7: use \hat{f}_{θ} to estimate optimal action sequence $\mathbf{A}_t^{(H)}$ (Eqn. 4)
- 8: execute first action a_t from selected action sequence $\mathbf{A}_t^{(H)}$
- 9: add (s_t, a_t) to \mathcal{D}_{RL}
- 10: **end for**
- 11: **end for**



1. Planning based on Model
2. Agent action selection

$$\hat{s}_{t+1} = s_t + \hat{f}_{\theta}(s_t, a_t)$$

$$\mathcal{E}(\theta) = \frac{1}{|\mathcal{D}|} \sum_{(s_t, a_t, s_{t+1}) \in \mathcal{D}} \frac{1}{2} \|(s_{t+1} - s_t) - \hat{f}_{\theta}(s_t, a_t)\|^2 \quad (1)$$

$$\mathbf{A}_t^{(H)} = \arg \max_{\mathbf{A}_t^{(H)}} \sum_{t'=t}^{t+H-1} r(\hat{s}_{t'}, a_{t'}) \quad :$$

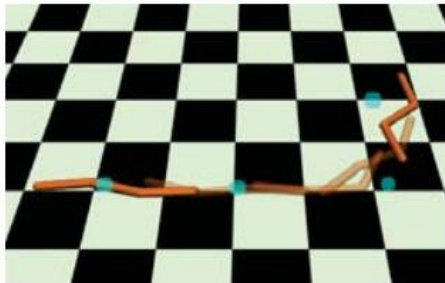
$$\hat{s}_t = s_t, \hat{s}_{t'+1} = \hat{s}_{t'} + \hat{f}_{\theta}(\hat{s}_{t'}, a_{t'}). \quad (3)$$

Random Shooting (RS):

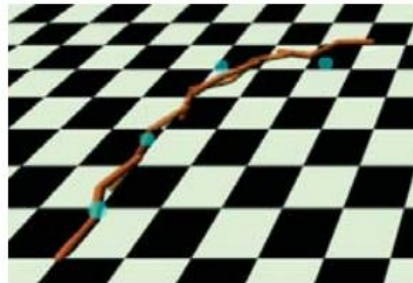
For $t = 1$ to T do:

$$\text{random}(a_t) \in A \rightarrow \mathcal{D}_{\text{RAND}}$$

□ MPC with Neural Network Model



(a) Swimmer left turn



(b) Swimmer right turn



(c) Ant left turn



(d) Ant right turn



(a) Swimmer



(b) Cheetah

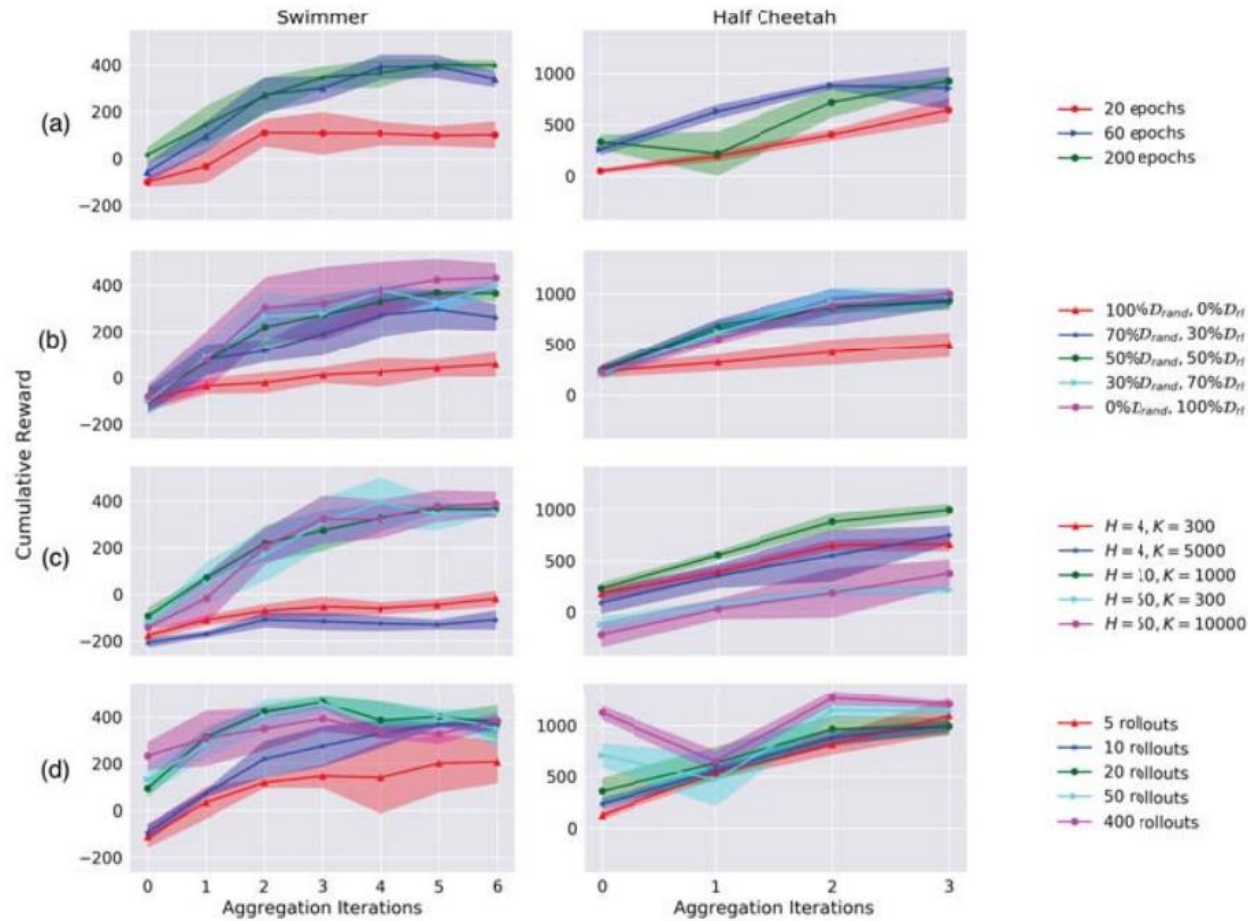


(c) Ant



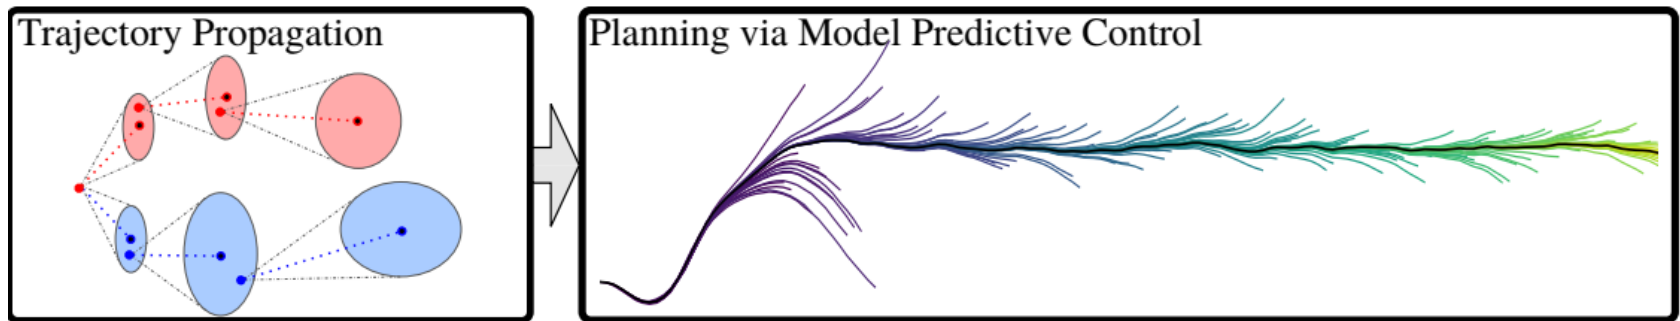
(d) Hopper

□ MPC with Neural Network Model

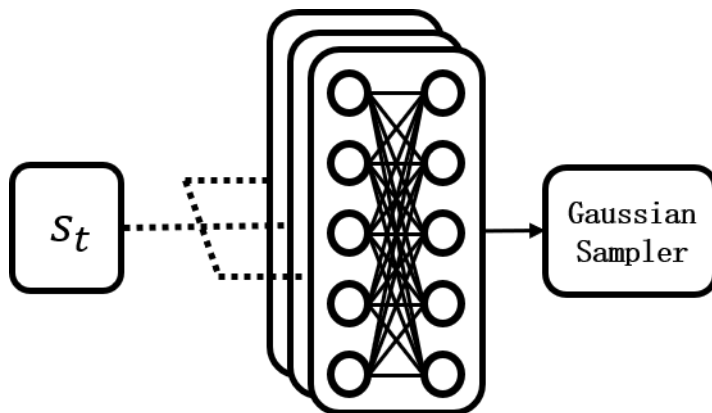


Decision-Time

□ Probabilistic Ensemble with Trajectory Sampling (PETS)



- Probabilistic Model Networks to capture **aleatoric uncertainty**.
- Model Ensembles to capture **epistemic uncertainty**.



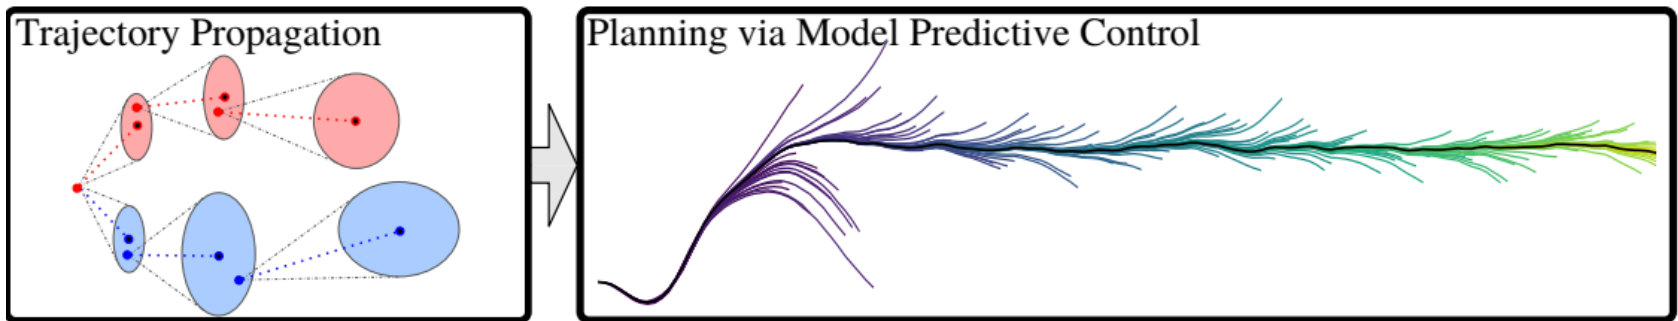
Model Ensemble Output:

$$\hat{s}_{t+1} = \text{random} (f_1(s_t, a_t), f_2(s_t, a_t), \dots, f_M(s_t, a_t))$$

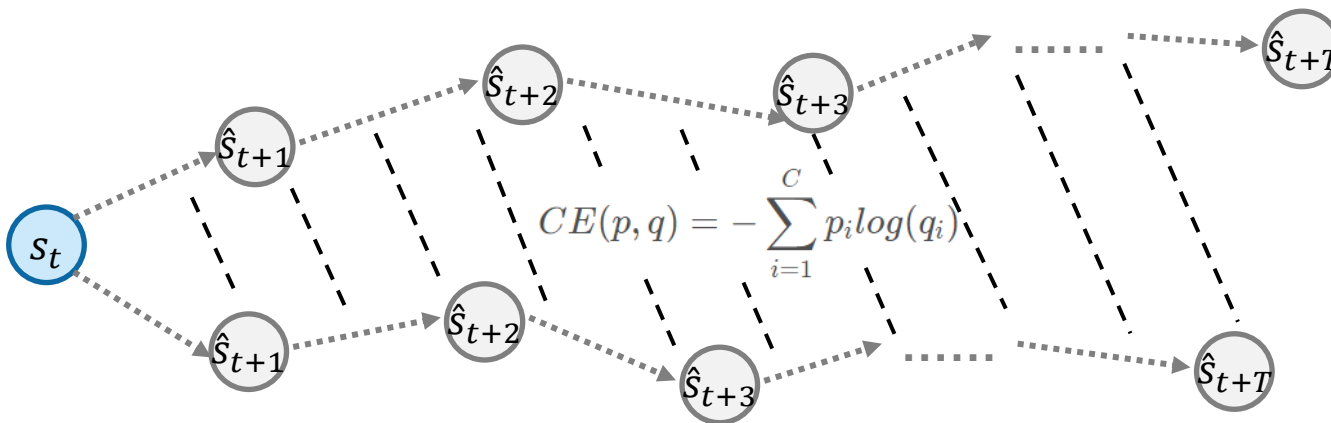
$$\hat{s}_{t+1} = \frac{1}{M} \sum_{j=1}^M f_j(s_t, a_t).$$

Decision-Time

□ Probabilistic Ensemble with Trajectory Sampling (PETS)

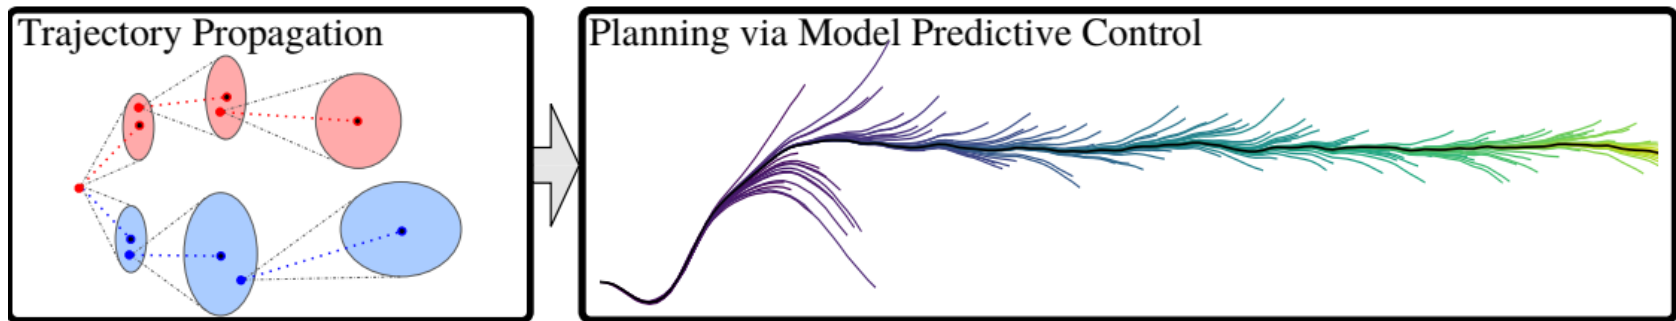


□ Cross Entropy Method (CEM) instead of Random Shooting (RS)



Decision-Time

Probabilistic Ensemble with Trajectory Sampling (PETS)

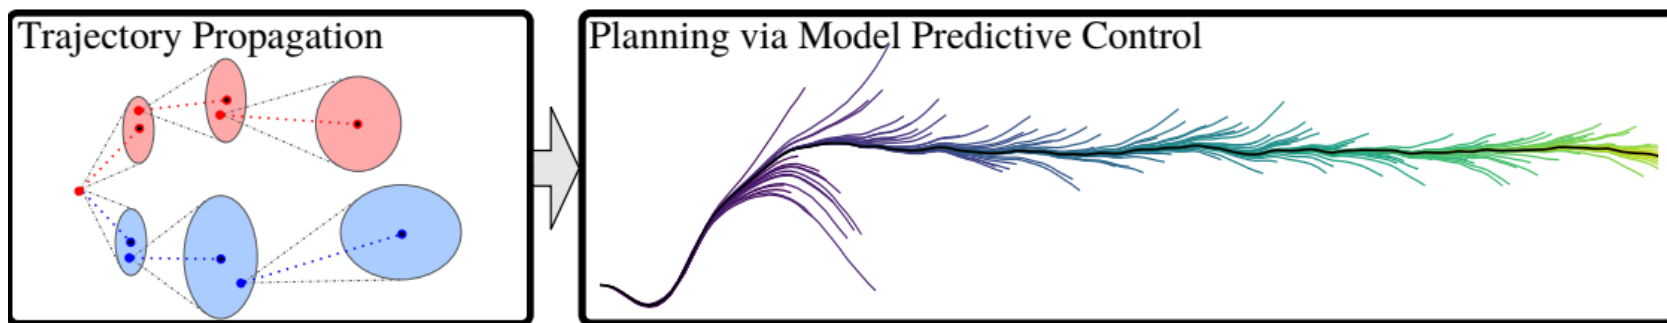


Algorithm 1 Our model-based MPC algorithm ‘PETS’:

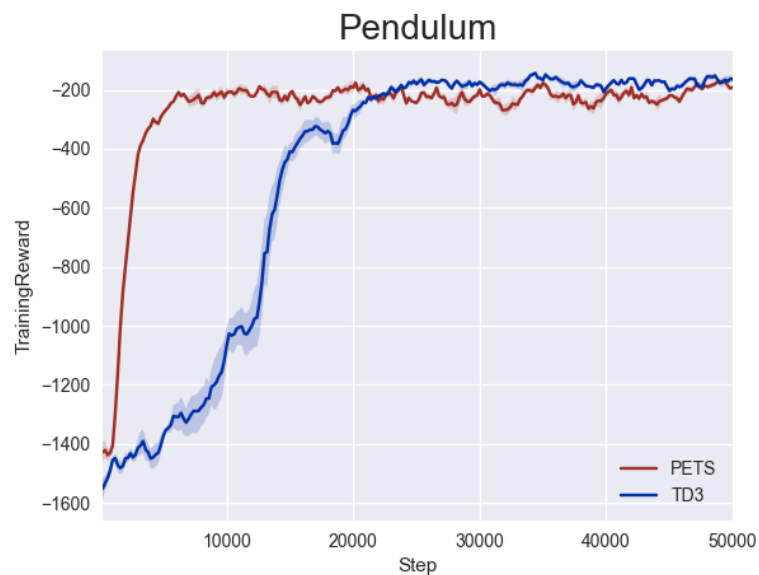
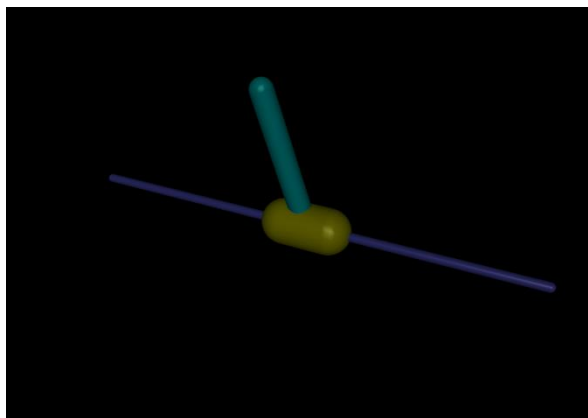
- 1: Initialize data \mathbb{D} with a random controller for one trial.
 - 2: **for** Trial $k = 1$ to K **do**
 - 3: Train a *PE* dynamics model \tilde{f} given \mathbb{D} .
 - 4: **for** Time $t = 0$ to TaskHorizon **do**
 - 5: **for** Actions sampled $\mathbf{a}_{t:t+T} \sim \text{CEM}(\cdot)$, 1 to $\tilde{N}\text{Samples}$ **do**
 - 6: Propagate state particles \mathbf{s}_τ^p using *TS* and $\tilde{f}|\{\mathbb{D}, \mathbf{a}_{t:t+T}\}$.
 - 7: Evaluate actions as $\sum_{\tau=t}^{t+T} \frac{1}{P} \sum_{p=1}^P r(\mathbf{s}_\tau^p, \mathbf{a}_\tau)$
 - 8: Update $\text{CEM}(\cdot)$ distribution.
 - 9: Execute first action \mathbf{a}_t^* (only) from optimal actions $\mathbf{a}_{t:t+T}^*$.
 - 10: Record outcome: $\mathbb{D} \leftarrow \mathbb{D} \cup \{\mathbf{s}_t, \mathbf{a}_t^*, \mathbf{s}_{t+1}\}$.
-

Decision-Time

□ Probabilistic Ensemble with Trajectory Sampling (PETS)



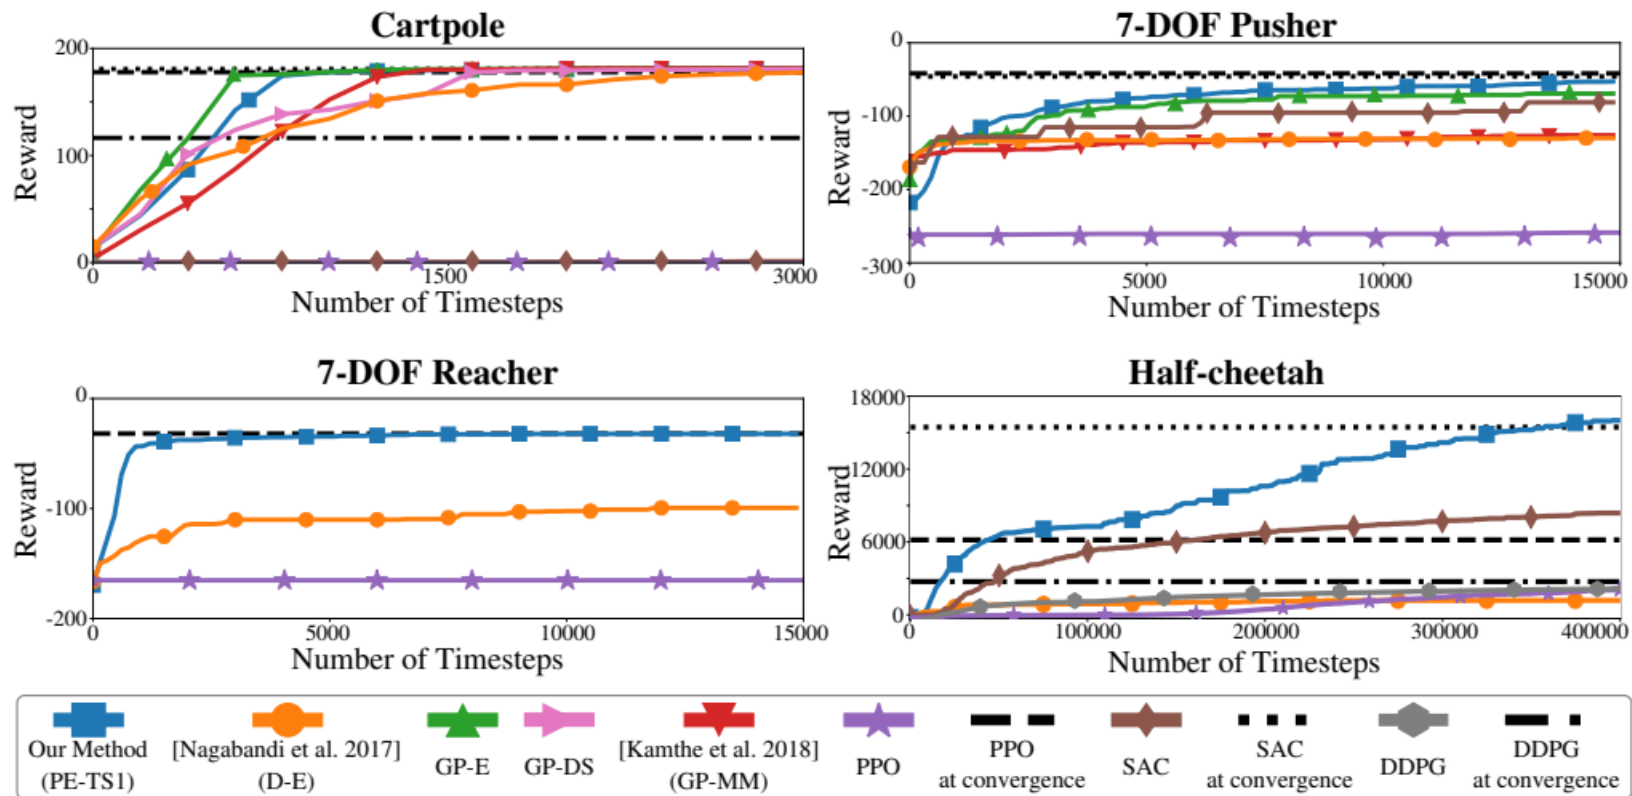
□ Experiments



Decision-Time

Probabilistic Ensemble with Trajectory Sampling (PETS)

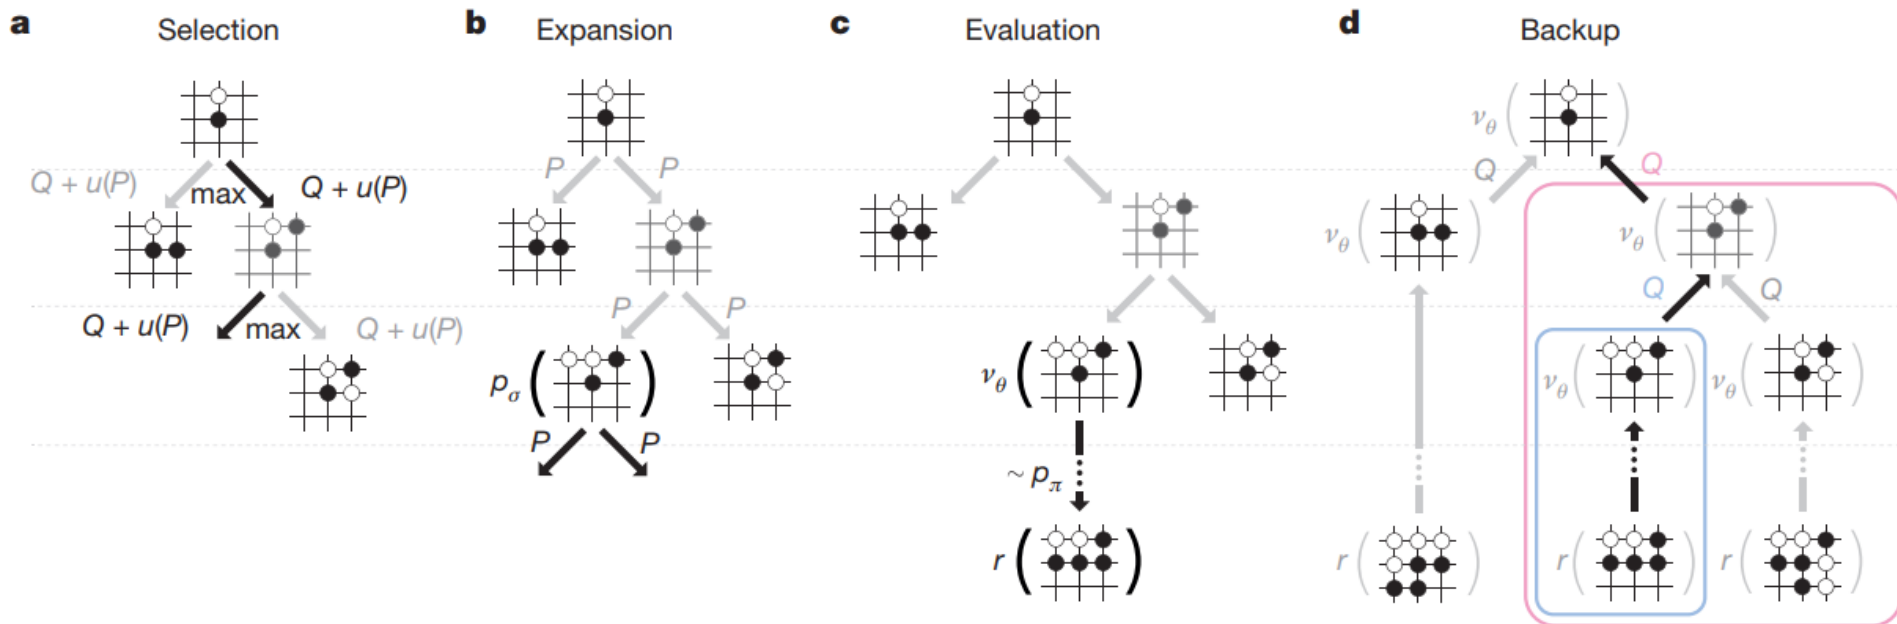
Experiments



Decision-Time

□ MuZero: Mastering Atari, Go and Chess

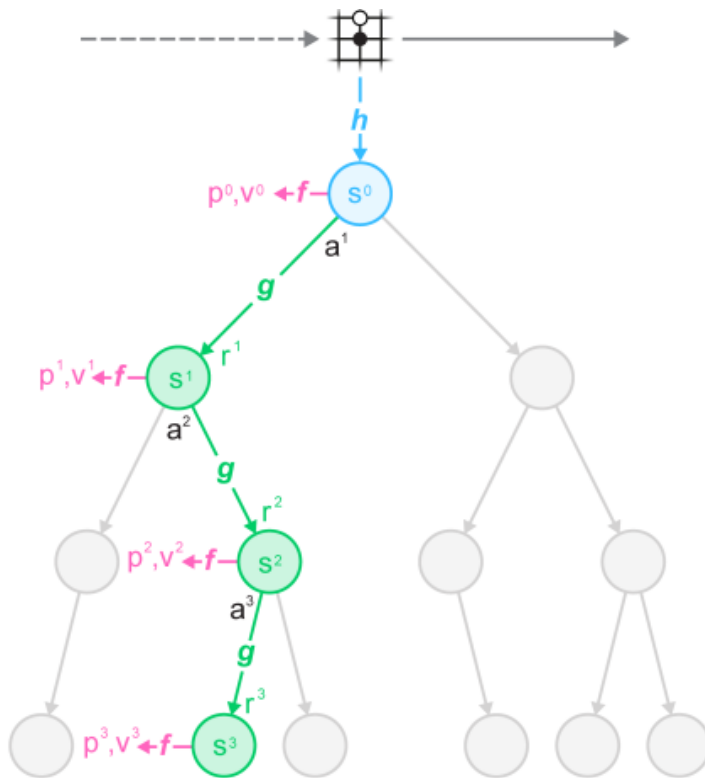
□ Monte-Carlo Tree Search (MCTS) with Dynamics Model



Decision-Time

□ MuZero: Mastering Atari, Go and Chess

□ Monte-Carlo Tree Search (MCTS) with Dynamics Model



□ Representation Function: $s^0 = h_{\theta}(o_1, \dots, o_n)$ to encode past observations.

□ Dynamic Function: $r^k, s^k = g_{\theta}(s^{k-1}, a^k)$

□ Prediction Function: $p^k, v^k = f_{\theta}(s^k)$

MCTS Backup:

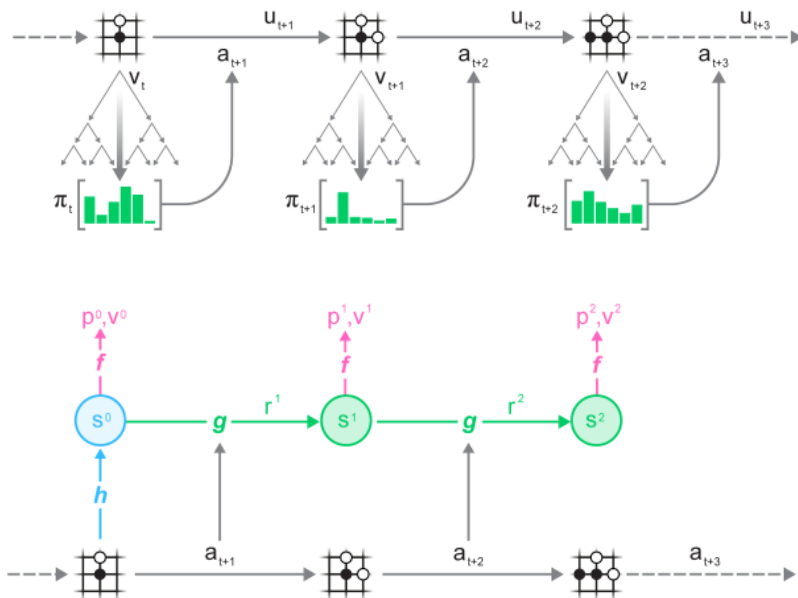
$$Q_{k+1}(s, a) = \frac{Q_k(s, a) \cdot N_k(s, a) + R}{N_k(s, a) + 1}$$

$$N_{k+1}(s, a) = N_k(s, a) + 1$$

Decision-Time

□ MuZero: Mastering Atari, Go and Chess

□ Monte-Carlo Tree Search (MCTS) with Dynamics Model



□ Representation Function: $s^0 = h_{\theta}(o_1, \dots, o_n)$ to encode past observations.

□ Dynamic Function: $r^k, s^k = g_{\theta}(s^{k-1}, a^k)$

□ Prediction Function: $p^k, v^k = f_{\theta}(s^k)$

MCTS Backup:

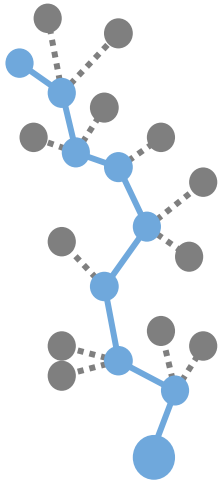
$$Q_{k+1}(s, a) = \frac{Q_k(s, a) \cdot N_k(s, a) + R}{N_k(s, a) + 1}$$

$$N_{k+1}(s, a) = N_k(s, a) + 1$$

Background

□ Dyna-Q

- Mix **real** and **model-generated** experiences and apply for additional policy updates.



Tabular Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

- $S \leftarrow$ current (nonterminal) state
- $A \leftarrow \epsilon$ -greedy(S, Q)
- Take action A ; observe resultant reward, R , and state, S'
- $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$
- $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)
- Loop repeat n times:

Update Q values with
real transitions

$S \leftarrow$ random previously observed state
 $A \leftarrow$ random action previously taken in S
 $R, S' \leftarrow Model(S, A)$

Update Q values with
model transitions

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?

Definition: Errors in the model can be exploited during policy optimization, resulting in large discrepancies between the predicted returns of the policy under the model and under the true dynamics.

$$\eta[\pi] \geq \hat{\eta}[\pi] - C.$$

$\eta[\pi]$: the returns of the policy in true MDP; $\hat{\eta}[\pi]$: the returns under learned model.

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?

Definition: Errors in the model can be exploited during policy optimization, resulting in large discrepancies between the predicted returns of the policy under the model and under the true dynamics.

$$\eta[\pi] \geq \hat{\eta}[\pi] - C.$$

The gap between true and model returns can be expressed in:

□ **Generalization Error:** $\epsilon_m = \mathbb{E}_{s,a \sim D_{env}} [D_{TV}(p(s'|s, a) || p_{\theta}(s'|s, a))]$

□ **Policy Distribution:** $\epsilon_{\pi} = D_{TV}(\pi_D(s) || \pi(s))$

Background

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?

The gap between:

$$\begin{aligned} |\eta_1 - \eta_2| &= \left| \sum_{s,a} (p_1(s,a) - p_2(s,a)) r(s,a) \right| \\ &= \left| \sum_{s,a} \left(\sum_t \gamma^t p_1^t(s,a) - p_2^t(s,a) \right) r(s,a) \right| \\ &= \left| \sum_t \sum_{s,a} \gamma^t (p_1^t(s,a) - p_2^t(s,a)) r(s,a) \right| \\ &\leq \sum_t \sum_{s,a} \gamma^t |p_1^t(s,a) - p_2^t(s,a)| r(s,a) \\ &\leq r_{\max} \sum_t \sum_{s,a} \gamma^t |p_1^t(s,a) - p_2^t(s,a)| \end{aligned}$$

Background

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?

The gap between:

$$|\eta_1 - \eta_2| \leq r_{\max} \sum_t \sum_{s,a} \gamma^t |p_1^t(s, a) - p_2^t(s, a)|$$

Lemma:

Lemma B.1 (TVD of Joint Distributions). Suppose we have two distributions $p_1(x, y) = p_1(x)p_1(y|x)$ and $p_2(x, y) = p_2(x)p_2(y|x)$. We can bound the total variation distance of the joint as:

$$D_{TV}(p_1(x, y) || p_2(x, y)) \leq D_{TV}(p_1(x) || p_2(x)) + E_{x \sim p_1}[D_{TV}(p_1(y|x) || p_2(y|x))]$$

Background

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?

The gap between:

$$|\eta_1 - \eta_2| \leq r_{\max} \sum_t \sum_{s,a} \gamma^t |p_1^t(s, a) - p_2^t(s, a)|$$

$$\epsilon_m = \max_t E_{s \sim \pi_{D,t}} [D_{TV}(p(s', r|s, a) \| p_\theta(s', r|s, a))]$$

Lemma:

$$\max_s D_{TV}(\pi \| \pi_D) \leq \epsilon_\pi$$

Lemma B.2 (Markov chain TVD bound, time-varying). Suppose the expected KL-divergence between two transition distributions is bounded as $\max_t E_{s \sim p_1^t(s)} D_{KL}(p_1(s'|s) \| p_2(s'|s)) \leq \delta$, and the initial state distributions are the same – $p_1^{t=0}(s) = p_2^{t=0}(s)$. Then the distance in the state marginal is bounded as:

$$D_{TV}(p_1^t(s) \| p_2^t(s)) \leq t\delta$$

$$\delta = \epsilon_m + \epsilon_\pi$$

Background

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?

The gap between:

$$\begin{aligned}
 |\eta_1 - \eta_2| &\leq r_{\max} \sum_t \sum_{s,a} \gamma^t |p_1^t(s,a) - p_2^t(s,a)| \begin{cases} \dashrightarrow D_{TV}(p_1^t(s)||p_2^t(s)) \leq t(\epsilon_m + \epsilon_\pi) \\ \dashrightarrow D_{TV}(\pi_1(a|s)||\pi_2(a|s)) \leq \epsilon_\pi \end{cases} \\
 &\leq 2r_{\max} \sum_t \gamma^t t(\epsilon_m + \epsilon_\pi) + \epsilon_\pi \\
 &\leq 2r_{\max} \left(\frac{\gamma(\epsilon_\pi + \epsilon_m)}{(1-\gamma)^2} + \frac{\epsilon_\pi}{1-\gamma} \right)
 \end{aligned}$$

The returns and model returns of the policy are bounded as:

$$\eta[\pi] \geq \hat{\eta}[\pi] - \underbrace{\left[\frac{2\gamma r_{\max}(\epsilon_m + 2\epsilon_\pi)}{(1-\gamma)^2} + \frac{4r_{\max}\epsilon_\pi}{(1-\gamma)} \right]}_{C(\epsilon_m, \epsilon_\pi)}$$

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?
- H-steps Q-target objective:

Model-based Critic Loss Function:

$$Loss_{critic} = \frac{1}{2} (y - Q(s_t, a_t))^2$$

Dyna-Q one-step Q-target:

$$y = \hat{r}(s_t, a_t) + \gamma Q(\hat{s}_{t+1}, \hat{a}_{t+1})$$

Multi-steps Q-target:

$$y = \sum_{k=t}^{H-1} \gamma^{k-t} \hat{r}_k + \gamma^H Q(\hat{s}_H, \hat{a}_H)$$

Background

□ Model-Based Policy Optimization (MBPO)

- Dyna-Q updates the policy with both real and model-generated transitions.
- How to guarantee the policy improvement if we only use the model-generated transitions ?
- H-steps Q-target objective
- MBPO Algorithm:

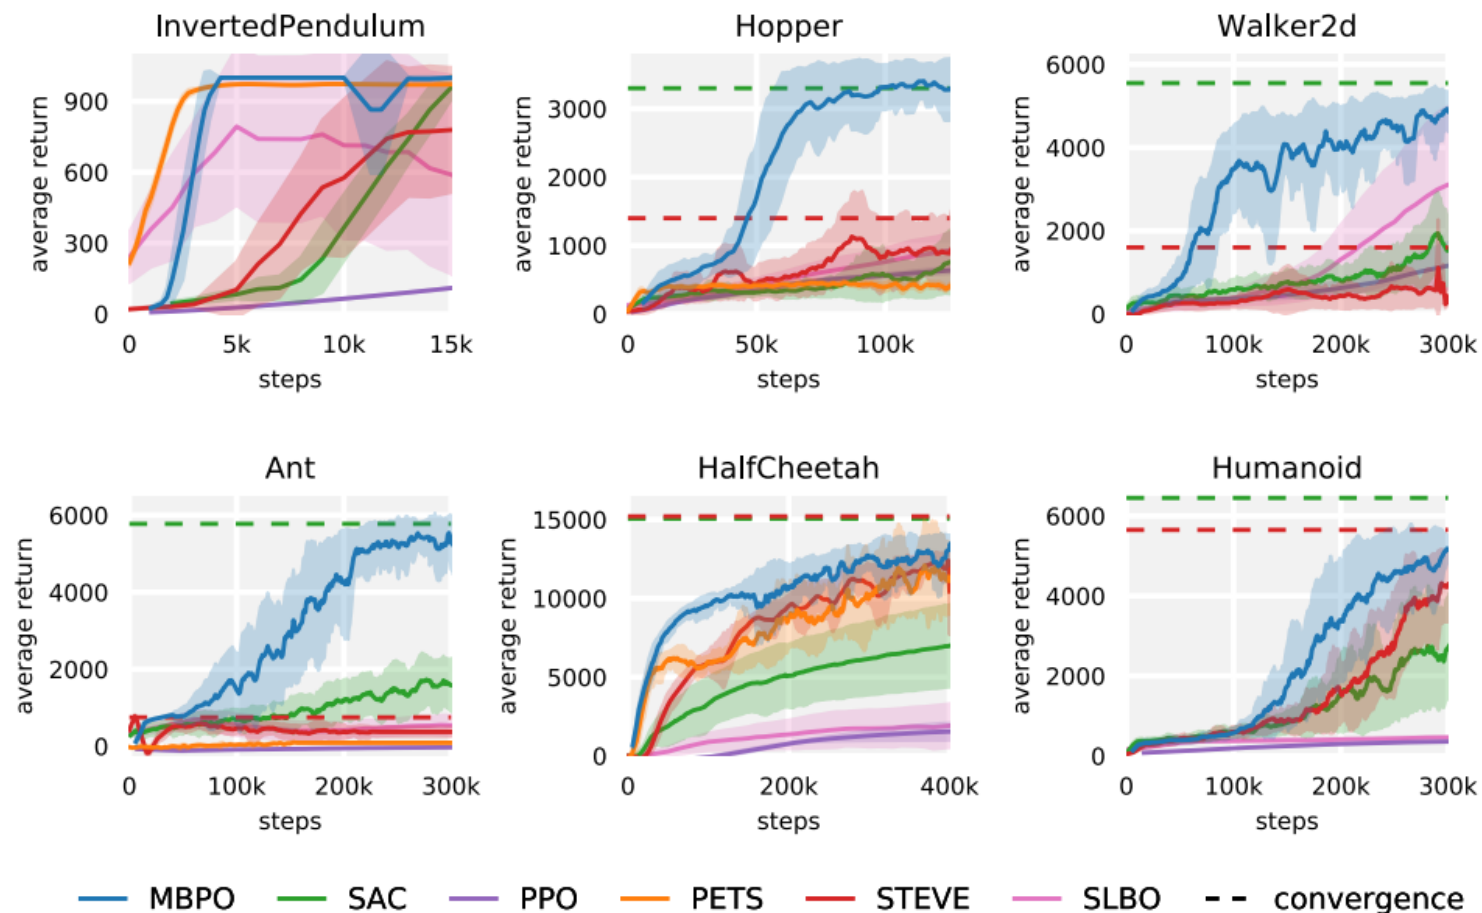
Algorithm 2 Model-Based Policy Optimization with Deep Reinforcement Learning

- 1: Initialize policy π_ϕ , predictive model p_θ , environment dataset \mathcal{D}_{env} , model dataset $\mathcal{D}_{\text{model}}$
 - 2: **for** N epochs **do**
 - 3: Train model p_θ on \mathcal{D}_{env} via maximum likelihood
 - 4: **for** E steps **do**
 - 5: Take action in environment according to π_ϕ ; add to \mathcal{D}_{env}
 - 6: **for** M model rollouts **do**
 - 7: Sample s_t uniformly from \mathcal{D}_{env}
 - 8: Perform k -step model rollout starting from s_t using policy π_ϕ ; add to $\mathcal{D}_{\text{model}}$
 - 9: **for** G gradient updates **do**
 - 10: Update policy parameters on model data: $\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi, \mathcal{D}_{\text{model}})$
-

Background

□ Model-Based Policy Optimization (MBPO)

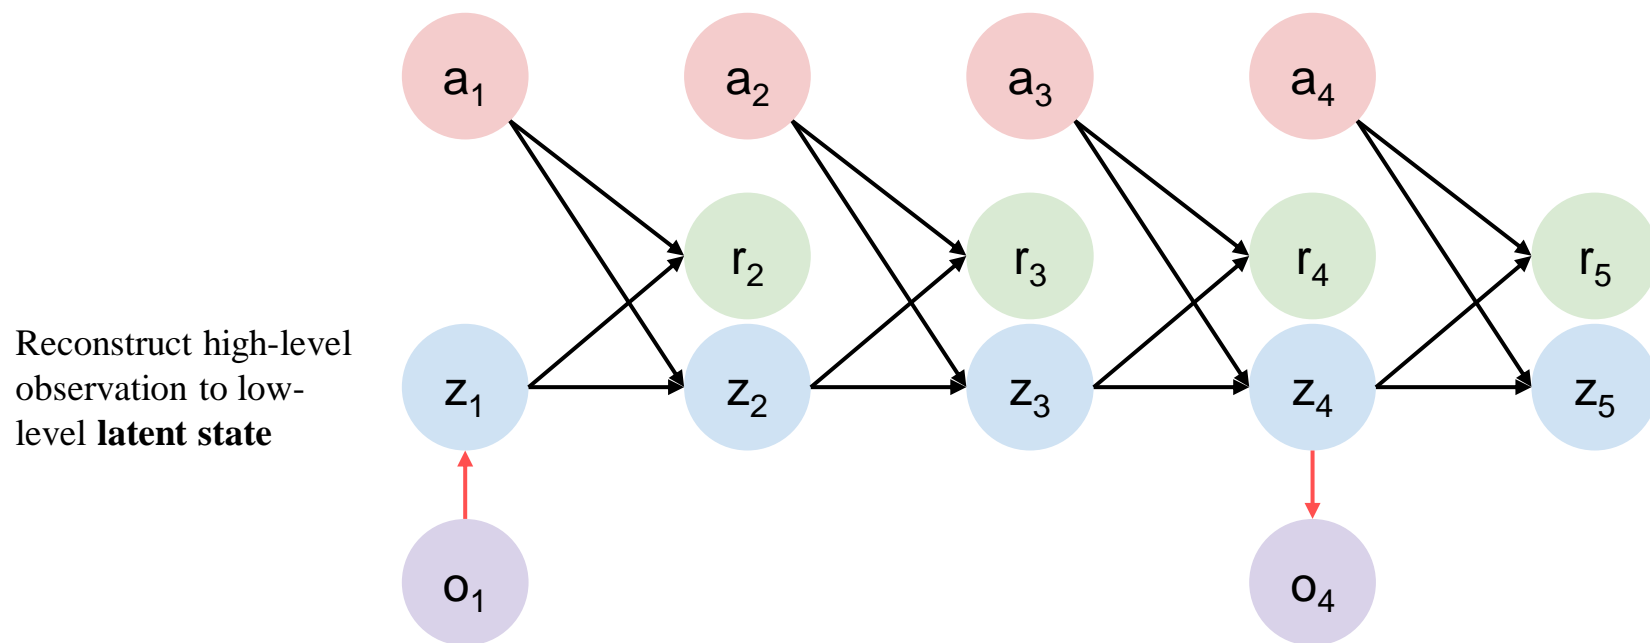
□ Experiments



Model-based RL with Images

□ Dreamer: Latent Imagination

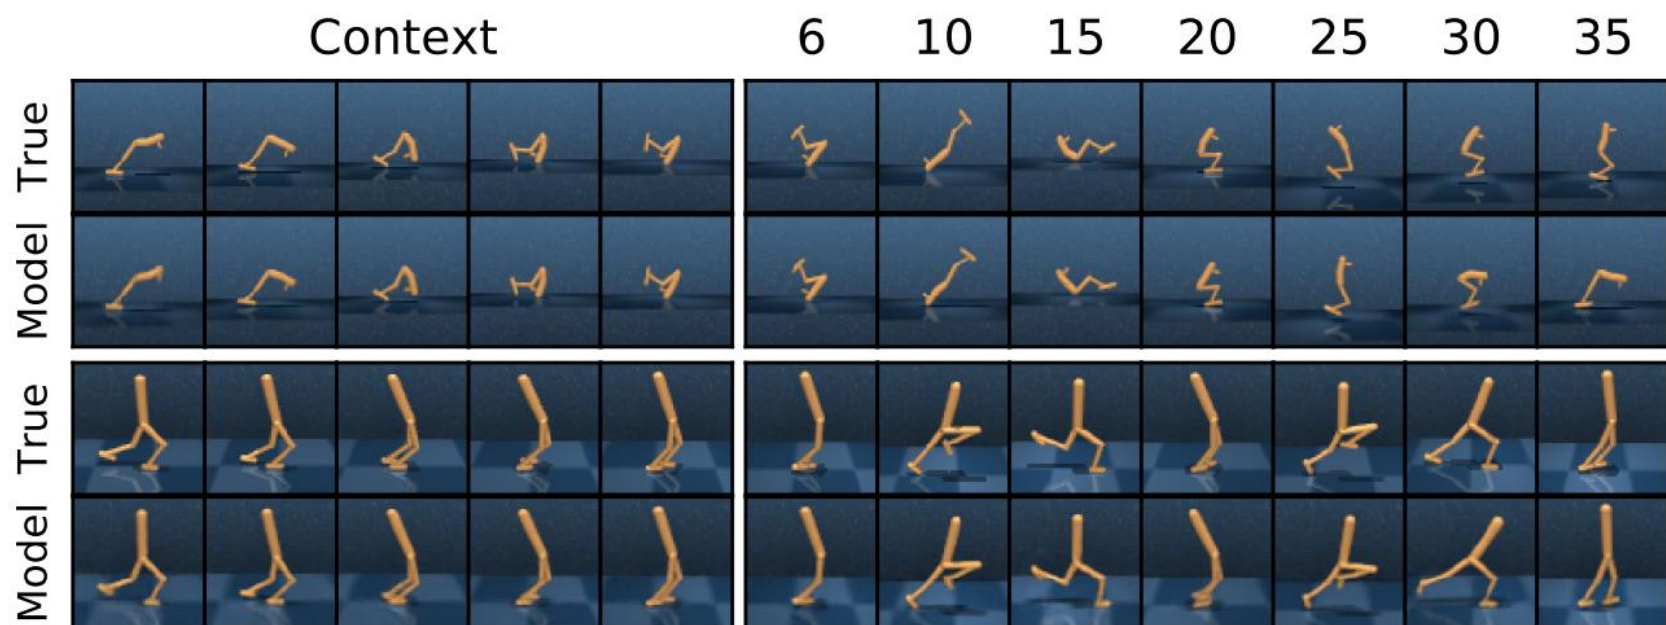
□ Latent state-transition models



Model-based RL with Images

□ Dreamer: Latent Imagination

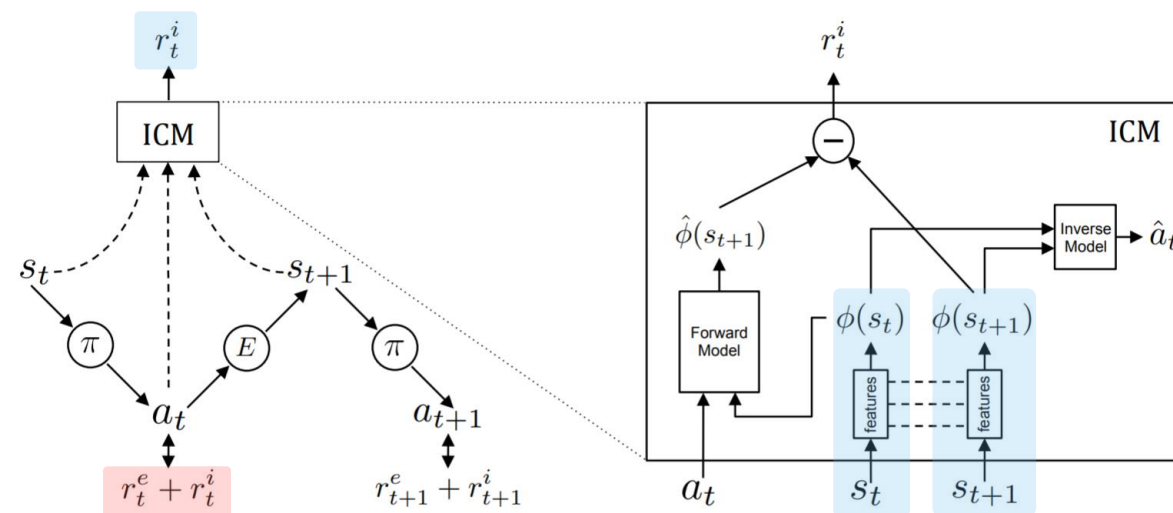
□ Latent state-transition models



Inverse Model

□ Intrinsic Curiosity Module (ICM)

- Besides extrinsic reward, agent set **intrinsic reward** to express its familiarity with the true environment. (Curiosity Driven)
- Model Prediction: $f(s_t, a_t) \rightarrow s_{t+1}$
- Intrinsic Reward: $r_t^{in} \sim \|f(s_t, a_t) - s_{t+1}\|_2^2$



- **Forward Dynamics Model:**

$$f_{\psi_F}(\phi(s_t), a_t) \rightarrow \hat{\phi}(s_{t+1})$$

- **Inverse Model:**

$$g_{\psi_I}(\phi(s_t), \phi(s_{t+1})) \rightarrow \hat{a}_t$$

- **Intrinsic Reward:**

$$r_t^i = \|\hat{\phi}(s_{t+1}) - \phi(s_{t+1})\|_2^2$$

Deep Reinforcement Learning



谢谢!

中山大学计算机学院